

Two Perspectives to Learning with Circuits



Motivation

Given a selection of sushi...



...and people's preferences...

Alice:     

Bob:     

Carol:     

...how can we model this as a probability distribution...

$$p(1^{\text{st}} = \text{salmon nigiri}, 3^{\text{rd}} = \text{salmon nigiri})$$

$$p(2^{\text{nd}} = \text{salmon nigiri} \mid 1^{\text{st}} = \text{maki roll})$$

$$\arg \max p(1^{\text{st}} = ?, 2^{\text{nd}} = ?, 3^{\text{rd}} = ?, 4^{\text{th}} = \text{maki roll}, 5^{\text{th}} = \text{tuna nigiri})$$

$$p((3^{\text{rd}} = \text{salmon nigiri} \rightarrow 1^{\text{st}} = \text{maki roll}) \vee 2^{\text{nd}} = \text{salmon nigiri})$$

...and extract meaningful queries from it?

Motivation

Given a selection of sushi...



...and people's preferences...

Alice:     

Bob:     

Carol:     

...how can we model this as a probability distribution...

$$p(1^{\text{st}} = \text{salmon nigiri}, 3^{\text{rd}} = \text{tuna nigiri})$$

$$p(2^{\text{nd}} = \text{tuna nigiri} \mid 1^{\text{st}} = \text{white rice ball})$$

$$\arg \max p(1^{\text{st}} = ?, 2^{\text{nd}} = ?, 3^{\text{rd}} = ?, 4^{\text{th}} = \text{white rice ball}, 5^{\text{th}} = \text{maki roll})$$

$$p((3^{\text{rd}} = \text{salmon nigiri} \rightarrow 1^{\text{st}} = \text{white rice ball}) \vee 2^{\text{nd}} = \text{tuna nigiri})$$

Marginals



Conditionals

MPE




Logical events

...and extract meaningful queries from it?


Probabilistic Circuits – Inputs



Probabilistic Circuits – Sums





Probabilistic Circuits – Smoothness




Definition 1 (Smoothness).


Every sum node child mentions the same variables.







Probabilistic Circuits – Determinism




Definition 2 (Determinism).
At most one sum node child has a positive value.





Probabilistic Circuits – Products






Probabilistic Circuits – Decomposability



Definition 3 (Decomposability).
Every product node child mentions different variables.



Probabilistic Circuits – Structured Decomposability



Definition 4 (Structured decomposability). *Every product node follows a vtree decomposition.*



Probabilistic Circuits – Tractability

Query	+Sm?	+Dec?	+Det?	+Str Dec?
Evidence	✓	✓	✓	✓
Marginals	✗	✓	✓	✓
Conditionals	✗	✓	✓	✓
MPE	✗	✗	✓	✓
Shannon Entropy	✗	✗	✓	✓
Rényi Entropy	✗	✗	✓	✓
Cross Entropy	✗	✗	✗	✓
Kullback-Leibler Div	✗	✗	✗	✓
Rényi's Alpha Div	✗	✗	✗	✓
Cauchy-Schwarz Div	✗	✗	✗	✓
Logical Events	✗	✗	✗	✓
Mutual Information	✗	✗	✗	✓

Probabilistic Circuits – Logic Circuits

A	B	C	$\phi(\mathbf{x})$
0	0	0	1
1	0	0	1
0	1	0	0
1	1	0	0
0	0	1	1
1	0	1	1
0	1	1	0
1	1	1	1



$$\phi(A, B, C) = (A \vee B) \wedge (\neg B \vee C)$$



Probabilistic Circuits – Support

A	B	C	$\phi(\mathbf{x})$	$p(\mathbf{x})$
0	0	0	1	0.140
1	0	0	1	0.024
0	1	0	0	0.000
1	1	0	0	0.000
0	0	1	1	0.560
1	0	1	1	0.096
0	1	1	0	0.000
1	1	1	1	0.180

$$\phi(A, B, C) = (A \vee B) \wedge (\neg B \vee C)$$




Learning Probabilistic Circuits – Where are we right now?

Name	Class	Time Complexity	# hyperparams	Accepts logic?	Sm?	Dec?	Det?	Str Dec?	$\{0, 1\}$?	\mathbb{N} ?	\mathbb{R} ?	Reference
LEARNSPN	DIV	$\begin{cases} \mathcal{O}(nkmc) & , \text{ if sum} \\ \mathcal{O}(nm^3) & , \text{ if product} \end{cases}$	≥ 2	✗	✓	✓	✗	✗	✓	✓	✓	Gens and Domingos [2013]
ID-SPN	DIV	$\begin{cases} \mathcal{O}(nkmc) & , \text{ if sum} \\ \mathcal{O}(nm^3) & , \text{ if product} \\ \mathcal{O}(ic(rn + m)) & , \text{ if input} \end{cases}$	$\geq 2 + 3$	✗	✓	✓	✗	✗	✓	✓	✗	Rooshenas and Lowd [2014]
PROMETHEUS	DIV	$\begin{cases} \mathcal{O}(nkmc) & , \text{ if sum} \\ \mathcal{O}(m(\log m)^2) & , \text{ if product} \end{cases}$	≥ 1	✗	✓	✓	✗	✗	✓	✓	✓	Jaini et al. [2018]
LEARNPSSD	INCR	$\begin{cases} \mathcal{O}(m^2) & , \text{ top-down vtree} \\ \mathcal{O}(m^4) & , \text{ bottom-up vtree} \\ \mathcal{O}(i \mathcal{C} ^2) & , \text{ circuit structure} \end{cases}$	1	✓	✓	✓	✓	✓	✓	✗	✗	Liang et al. [2017]
STRUDEL	INCR	$\begin{cases} \mathcal{O}(m^2n) & , \text{ CLT + vtree} \\ \mathcal{O}(i(\mathcal{C} n + m^2)) & , \text{ circuit structure} \end{cases}$	1	✓	✓	✓	✓	✓	✓	✗	✗	Dang et al. [2020]
RAT-SPN	RAND	$\mathcal{O}(rd(s + l))$	4	✗	✓	✓	✗	✗	✓	✓	✓	Peharz et al. [2020]
XPC	RAND	$\mathcal{O}(i(t + kn) + ikm^2n)$	3	✗	✓	✓	✓	✓	✓	✗	✗	Mauro et al. [2021]
SAMPLEPSDD	RAND	$\begin{cases} \mathcal{O}(m) & , \text{ random vtree} \\ \mathcal{O}(kc \log c + \log_2^2 k) & , \text{ per call} \end{cases}$	1	✓	✓	✓	✓	✓	✓	✗	✗	Geh and Mauá [2021]
LEARNRP	RAND	$\begin{cases} \mathcal{O}(m^2) & , \text{ top-down vtree} \\ \mathcal{O}(m^4) & , \text{ bottom-up vtree} \\ \mathcal{O}(knm) & , \text{ per call} \end{cases}$	0	✗	✓	✓	✗	✓	✓	✓	✓	To appear

A Logical Perspective

Motivation



Example:

$$n = 3, k = 3$$

X_{11}	X_{12}	X_{13}	X_{21}	\dots	X_{33}	$p(\mathbf{x}) > 0$
0	0	0	0	0	0	0
1	0	0	0	0	0	0
0	1	0	0	0	0	0
1	1	0	0	0	0	0
0	0	1	0	0	0	0
\vdots	\vdots	\vdots	\vdots	\vdots	\vdots	\vdots
0	1	1	1	1	1	0
1	1	1	1	1	1	0

Assignments: $2^{3 \cdot 3} = 512$

Positive assignments: $3! = 6$

If we assume

n sushi types,

k sized rankings with $k \leq n$,

X_{ij} binary variables; i is sushi type, j is position in ranking;

then the total number of possible assignments of the $n \cdot k$ variables is 2^{nk} ...

...but many of these are zero probability assignments!

If we can embed total ranking constraints...

...we go down to $k!$ total assignments!

Takeaway: models which exploit domain knowledge are much more efficient!

Motivation

Existing approaches:

LEARNPSDD (Liang et al. [2017]):


- ✗ Requires initial logic circuit encoding the support...
- ✗ Scales poorly to complex formulae and/or high dimension...
- ✗ Costly whole circuit evaluation at every iteration...
- ✓ Very good performance!

STRUDEL (Dang et al. [2020]):

- ✓ Constructs an initial structure (from a CLT)!
- ✗ But does not encode constraints...
- ✓ Scales to high dimension!
- ✗ As long as the circuit doesn't get too big...

SAMPLEPSDD (Geh and Mauá [2021]):



- ✓ Scales to high dimension and complex formulae!
- ✓ Constructs a structure consistent with constraints!
- ✗ But does so by relaxing the formula...
- ✗ Performance varies on set bounds and vtree structure...



SAMPLEPSDD

Common assumption: p_i are conjunctions of literals.

$$\phi(A, B, C, D) = (A \wedge \neg B \wedge \neg D) \vee (B \wedge \neg C \wedge D)$$





Problem: size of circuit is **exponential** in the size of p_i 's scope.

SAMPLEPSDD

Solution: randomly sample a bounded number (k) of p_i

$$\phi(A, B, C, D) = (A \wedge \neg B \wedge \neg D) \vee (B \wedge \neg C \wedge D)$$



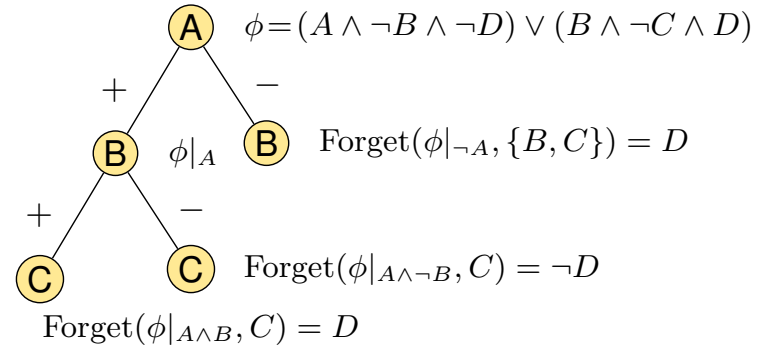

But: this violates structured decomposability:

$\neg C \wedge D$ contains C , and $C \notin S$


$\neg B \wedge \neg C \wedge D$ contains B and C , and $B, C \notin S$

SAMPLEPSDD

New solution: relax logical constraints ϕ




Now all s_i respect S



SAMPLEPSDD

Apply **local transformations** for variety and size reduction



Experiments

Evaluation: we sample 30 PSDDs and use 5 ensemble strategies:

- Likelihood weighting (LLW),
- Uniform weights,
- ◆ Expectation Maximization (EM),
- ▲ Stacking,
- ▼ Bayesian Model Combination (BMC);



comparing against **STRUDEL**, **LEARNPSDD** and **LEARNSPN**.

Datasets: we evaluate with 5 data + knowledge as logic constraints:

	Dataset	#vars	#train	ϕ 's size
⇒	LED	14	5000	23
⇒	LED + IMAGES	157	700	39899
	SUSHI RANKING	100	3500	17413
	SUSHI TOP 5	10	3500	37
	DOTA 2 GAMES	227	92650	1308

Our approach fares **better** with **fewer** data, yet
remains **competitive** under **lots** of data.

Mattei et al. [2020], Kamishima [2003], Shen et al. [2017],
Choi et al. [2015], Gens and Domingos [2013], Dang et al. [2020]



Experiments

Evaluation: we sample 30 PSDDs and use 5 ensemble strategies:

- Likelihood weighting (LLW),
- Uniform weights,
- ◆ Expectation Maximization (EM),
- ▲ Stacking,
- ▼ Bayesian Model Combination (BMC);


comparing against **STRUDEL**, **LEARNPSDD** and **LEARNSPN**.

Datasets: we evaluate with 5 data + knowledge as logic constraints:

Dataset	#vars	#train	ϕ 's size
LED	14	5000	23
LED + IMAGES	157	700	39899
⇒ SUSHI RANKING	100	3500	17413
⇒ SUSHI TOP 5	10	3500	37
DOTA 2 GAMES	227	92650	1308

Our approach fares **better** with **fewer** data, yet
remains **competitive** under **lots of data**.

Mattei et al. [2020], Kamishima [2003], Shen et al. [2017],
Choi et al. [2015], Gens and Domingos [2013], Dang et al. [2020]



Experiments

Evaluation: we sample 30 PSDDs and use 5 ensemble strategies:

- Likelihood weighting (LLW),
- Uniform weights,
- ◆ Expectation Maximization (EM),
- ▲ Stacking,
- ▼ Bayesian Model Combination (BMC);



comparing against **STRUDEL**, **LEARNPSDD** and **LEARNSPN**.

Datasets: we evaluate with 5 data + knowledge as logic constraints:

Dataset	#vars	#train	ϕ 's size
LED	14	5000	23
LED + IMAGES	157	700	39899
SUSHI RANKING	100	3500	17413
SUSHI TOP 5	10	3500	37
⇒ DOTA 2 GAMES	227	92650	1308



Our approach fares **better** with **fewer** data, yet
remains **competitive** under **lots of data**.

Mattei et al. [2020], Kamishima [2003], Shen et al. [2017],
Choi et al. [2015], Gens and Domingos [2013], Dang et al. [2020]





Experiments

What is the impact of higher k 's and right-leaning vtrees
in log-likelihood and consistency?



Samples perform better with higher k 's and right-leaning vtrees ...
...but at a cost to complexity.




What do we gain from this?

Available queries:

- ✓ Probability of Evidence;
- ✓ Marginal Probability;
- ✓ Conditional Probability;
- ✓ Most Probable Explanation;
- ✓ Shannon Entropy;
- ✓ Cross Entropy;
- ✓ Kullback-Leibler Divergence;
- ✓ Rényi's Alpha Divergence;
- ✓ Cauchy-Schwarz Divergence;
- ✓ Probability of Logical Events;
- ✓ Mutual Information.


A	B	C	$p(\mathbf{x})$
0	0	0	0.1
0	1	0	0.1
1	0	0	0.2
1	0	1	0.6

$$\phi(A, B, C) = (A \rightarrow \neg B) \wedge (C \rightarrow A)$$



Support:

- ✓ Defineable as a logic formula;
- ✗ Consistent with a relaxation;
- ✓ Ensembles mitigate relaxation.



A Data Perspective

Motivation




Density Estimation Trees...

- ✓ ...are fast;
- ✓ ...are interpretable;
- ✓ ...are (somewhat) explainable;
- ✓ ...have extensive literature coverage;
- ✗ ...are not so expressive;
- ✗ ...only accept marginalization queries;
- ✗ ...are not so accurate;


...but are subsumed by circuits!

Learn DETs \subseteq Learn PCs?


Can we take advantage of known learning procedures in DETs and transplant them to more general circuits?



Random Projections





Axis-aligned projections



Random projections



If the data has *intrinsic dimension* d , then with constant probability the part of the data at level d or higher of the tree has average diameter less than half of the data.

Random Projections






If the data has *intrinsic dimension* d , then with constant probability the part of the data at level d or higher of the tree has average diameter less than half of the data.

LearnRP



LearnRP



LearnRP

